

# VariantAnnotation

Valerie Obenchain \*

Fred Hutchinson Cancer Research Center  
Seattle, WA

24 July 2012

## Abstract

This lab provides an overview of the [VariantAnnotation](#) package. Variants from a VCF file are read into R and stored in VCF-class object. We explore the class and see how the data are organized and accessed. Variants are then classified by region such as 'intronic', 'coding', '3UTR' etc. Coding variants are further analyzed for amino acid coding changes. Predictions as to how damaging these coding changes may be are determined by querying the SIFT database package.

## 1 Annotation of Variants

A major product of DNASeq experiments are catalogs of called variants (e.g., SNPs, indels). We will use the [VariantAnnotation](#) package to explore this type of data. Sample data included in the package are a subset of chromosome 22 from the [1000 Genomes](#) project. Variant Call Format (VCF; [full description](#)) text files contain meta-information lines, a header line with column names, data lines with information about a position in the genome, and optional genotype information on samples for each position.

### 1.1 Variant call format (VCF) files

Data are read from a VCF file and variants identified according to region such as coding, intron, intergenic, spliceSite etc. Amino acid coding changes are computed for the non-synonymous variants. SIFT and PolyPhen databases provide predictions of how severely the coding changes affect protein function.

#### Data exploration

##### Exercise 1

*The objective of this exercise is to compare the quality of called SNPs that are located in dbSNP, versus those that are novel.*

---

\*[vobencha@fhcrc.org](mailto:vobencha@fhcrc.org)

Locate the sample data in the file system. Explore the metadata (information about the content of the file) using `scanVcfHeader`. Discover the 'info' fields VT (variant type), and RSQ (genotype imputation quality).

Input sample data in using `readVcf`. You'll need to specify the genome build (`genome="hg19"`) on which the variants are annotated. Take a peak at the `rowData` to see the genomic locations of each variant.

`dbSNP` uses abbreviations such as `ch22` to represent chromosome 22, whereas the VCF file uses 22. Use `rowData` and `renameSeqlevels` to extract the row data of the variants, and rename the chromosomes.

The [SNPlocs.Hsapiens.dbSNP.20101109](#) contains information about SNPs in a particular build of `dbSNP`. Load the package, use the `dbSNPFilter` function to create a filter, and query the row data of the VCF file for membership.

Create a data frame containing the `dbSNP` membership status and imputation quality of each SNP. Create a density plot to illustrate the results.

**Solution:** Explore the header:

```
> library(VariantAnnotation)
> fl <- system.file("extdata", "chr22.vcf.gz",
+                   package="VariantAnnotation")
> (hdr <- scanVcfHeader(fl))
```

```
class: VCFHeader
samples(5): HG00096 HG00097 HG00099 HG00100 HG00101
meta(1): fileformat
fixed(1): ALT
info(22): LDAF AVGPOST ... VT SNPSOURCE
geno(3): GT DS GL
```

```
> info(hdr)[c("VT", "RSQ"),]
```

DataFrame with 2 rows and 3 columns

	Number	Type	Description
	<character>	<character>	<character>
VT	1	String	indicates what type of variant the line represents
RSQ	1	Float	Genotype imputation quality from MaCH/Thunder

Input the data and peak at their locations:

```
> vcf <- readVcf(fl, "hg19")
> head(rowData(vcf), 3)
```

GRanges with 3 ranges and 1 elementMetadata col:

seqnames	ranges	strand	paramRangeID
<Rle>	<IRanges>	<Rle>	<factor>
rs7410291	22 [50300078, 50300078]	*	<NA>
rs147922003	22 [50300086, 50300086]	*	<NA>
rs114143073	22 [50300101, 50300101]	*	<NA>

```

---
seqlengths:
  22
  NA

```

Rename chromosome levels:

```

> rowData(vcf) <- renameSeqlevels(rowData(vcf), c("22"="ch22"))

```

Discover whether SNPs are located in dbSNP:

```

> library(SNPlocs.Hsapiens.dbSNP.20101109)
> snpFilt <- dbSNPFilter("SNPlocs.Hsapiens.dbSNP.20101109")
> inDbSNP <- snpFilt(rowData(vcf), subset=FALSE)
> table(inDbSNP)

inDbSNP
FALSE  TRUE
 6126  4250

```

Create a data frame summarizing SNP quality and dbSNP membership:

```

> metrics <-
+   data.frame(inDbSNP=inDbSNP, RSQ=values(info(vcf))$RSQ)

```

Finally, visualize the data, e.g., using `ggplot2` (Figure 1).

```

> library(ggplot2)
> ggplot(metrics, aes(RSQ, fill=inDbSNP)) +
+   geom_density(alpha=0.5) +
+   scale_x_continuous(name="MaCH / Thunder Imputation Quality") +
+   scale_y_continuous(name="Density") +
+   opts(legend.position="top")

```

## 1.2 Coding consequences

**Locating variants in and around genes** Variant location with respect to genes can be identified with the `locateVariants` function. Regions are specified in the `region` argument and can be one of the following constructors: `CodingVariants()`, `IntronVariants()`, `FiveUTRVariants()`, `ThreeUTRVariants()`, `IntergenicVariants()`, `SpliceSiteVariants()`, or `AllVariants()`. Location definitions are shown in Table 1.

### Exercise 2

Load the *[TxDb.Hsapiens.UCSC.hg19.ensGene](#)* annotation with the `loadDb` function from *[GenomicFeatures](#)*. The annotation file is located at `"/home/valerie/VariantAnnotation/"`. Read in the `chr22.vcf.gz` example file from the *[VariantAnnotation](#)* package.

Remembering to re-name sequence levels, use the `locateVariants` function to identify coding variants.

Summarize aspects of your data, e.g., did any coding variants match more than one gene? How many coding variants are there per gene ID?




Figure 1: Quality scores of variants in dbSNP, compared to those not in dbSNP.

Table 1: Variant locations

Location	Details
coding	Within a coding region
fiveUTR	Within a 5' untranslated region
threeUTR	Within a 3' untranslated region
intron	Within an intron region
intergenic	Not within a transcript associated with a gene
spliceSite	Overlaps any of the first or last 2 nucleotides of an intron

**Solution:** The TxDb.Hsapiens.UCSC.hg19.ensGene.sqlite file contains annotations for the Ensembl gene model. Load the TranscripDb,

```
> library(GenomicFeatures) # for loadDb
> txdb <-
+   loadDb("/home/valerie/VariantAnnotation/TxDb.Hsapiens.UCSC.hg19.ensGene.sqlite")

and read in the VCF file.
```

```
> fl <- system.file("extdata", "chr22.vcf.gz",
+                   package="VariantAnnotation")
> vcf <- readVcf(fl, "hg19")
> vcf <- renameSeqlevels(vcf, c("22"="chr22"))
```

The next lines locate coding variants.

```
> rd <- rowData(vcf)
> loc <- locateVariants(rd, txdb, CodingVariants())
> head(loc, 3)
```

GRanges with 3 ranges and 5 elementMetadata cols:

	seqnames	ranges	strand	LOCATION	QUERYID	TXID
	<Rle>	<IRanges>	<Rle>	<factor>	<integer>	<integer>
[1]	chr22	[50301422, 50301422]	*	coding	24	165767
[2]	chr22	[50301422, 50301422]	*	coding	24	165768
[3]	chr22	[50301422, 50301422]	*	coding	24	165769

	CDSID	GENEID
	<integer>	<character>
[1]	256696	ENSG00000182858
[2]	256695	ENSG00000182858
[3]	256696	ENSG00000182858

---

```
seqlengths:
chr22
NA
```

To answer gene-centric questions data can be summarized by gene regardless of transcript.

```
> ## Did any coding variants match more than one gene?
> splt <- split(values(loc)$GENEID, values(loc)$QUERYID)
> table(sapply(splt, function(x) length(unique(x)) > 1))
```

```
FALSE
1026
```

```
> ## Summarize the number of coding variants by gene ID
> splt <- split(values(loc)$QUERYID, values(loc)$GENEID)
> head(sapply(splt, function(x) length(unique(x))), 3)
```

```

ENSG00000025708  ENSG00000025770  ENSG00000073146
      13              48              82

```

**Amino acid coding changes** `predictCoding` computes amino acid coding changes for non-synonymous variants. Only ranges in query that overlap with a coding region in subject are considered. Reference sequences are retrieved from either a `BSgenome` or fasta file specified in `seqSource`. Variant sequences are constructed by substituting, inserting or deleting values in the `varAllele` column into the reference sequence. Amino acid codes are computed for the variant codon sequence when the length is a multiple of 3.

The query argument to `predictCoding` can be a `GRanges` or `VCF`. When a `GRanges` is supplied the `varAllele` argument must be specified. In the case of a `VCF`, the alternate alleles are taken from `values(alt(<VCF>))$ALT` and the `varAllele` argument is not specified.

The result is a modified query containing only variants that fall within coding regions. Each row represents a variant-transcript match so more than one row per original variant is possible.

```

> library(BSgenome.Hsapiens.UCSC.hg19)
> coding <- predictCoding(vcf, txdb, seqSource=Hsapiens)
> coding[5:9]

```

`GRanges` with 5 ranges and 13 `elementMetadata` cols:

	seqnames	ranges	strand	paramRangeID
	<Rle>	<IRanges>	<Rle>	<factor>
rs8135963	chr22	[50301476, 50301476]	-	<NA>
rs8135963.1	chr22	[50301476, 50301476]	-	<NA>
22:50301488	chr22	[50301488, 50301488]	-	<NA>
22:50301488.1	chr22	[50301488, 50301488]	-	<NA>
22:50301488.2	chr22	[50301488, 50301488]	-	<NA>

	varAllele	CDSLOC	PROTEINLOC	QUERYID
	<DNAStringSet>	<IRanges>	<CompressedIntegerList>	<integer>
rs8135963	G	[ 91, 91]	31	25
rs8135963.1	G	[416, 416]	139	25
22:50301488	A	[873, 873]	291	26
22:50301488.1	A	[ 79, 79]	27	26
22:50301488.2	A	[404, 404]	135	26

	TXID	CDSID	GENEID	CONSEQUENCE
	<character>	<integer>	<character>	<factor>
rs8135963	165768	256695	ENSG00000182858	nonsynonymous
rs8135963.1	165769	256696	ENSG00000182858	nonsynonymous
22:50301488	165767	256696	ENSG00000182858	synonymous
22:50301488.1	165768	256695	ENSG00000182858	nonsynonymous
22:50301488.2	165769	256696	ENSG00000182858	nonsynonymous

	REFCODON	VARCODON	REFAA	VARAA
	<DNAStringSet>	<DNAStringSet>	<AAStringSet>	<AAStringSet>
rs8135963	ACT	GCT	T	A
rs8135963.1	CAC	CGC	H	R
22:50301488	CCG	CCA	P	P
22:50301488.1	GAC	AAC	D	N
22:50301488.2	CGA	CAA	R	Q

---

```

seqlengths:
chr22
NA

```

Using variant rs114264124 as an example, we see `varAllele` A has been substituted into the `refCodon` CGG to produce `varCodon` CAG. The `refCodon` is the sequence of codons necessary to make the variant allele substitution and therefore often includes more nucleotides than indicated in the range (i.e. the range is 50302962, 50302962, width of 1). Notice it is the second position in the `refCodon` that has been substituted. This position in the codon, the position of substitution, corresponds to genomic position 50302962. This genomic position maps to position 698 in coding region-based coordinates and to triplet 233 in the protein. This is a non-synonymous coding variant where the amino acid has changed from R (Arg) to Q (Gln).

When the resulting `varCodon` is not a multiple of 3 it cannot be translated. The consequence is considered a `frameshift` and `varAA` will be missing. There are no frameshifts in this file but we do have some variants classified as `nonsense` which indicates premature stop codons.

```
> table(values(coding)$CONSEQUENCE)
```

nonsense	nonsynonymous	synonymous
22	1884	1449

**SIFT and PolyPhen databases** From `predictCoding` we identified the amino acid coding changes for the non-synonymous variants. For this subset we can retrieve predictions of how damaging these coding changes may be. SIFT (Sorting Intolerant From Tolerant) and PolyPhen (Polymorphism Phenotyping) are methods that predict the impact of amino acid substitution on a human protein. The SIFT method uses sequence homology and the physical properties of amino acids to make predictions about protein function. PolyPhen uses sequence-based features and structural information characterizing the substitution to make predictions about the structure and function of the protein.

Collated predictions for specific dbSNP builds are available as downloads from the SIFT and PolyPhen web sites. These results have been packaged into *SIFT.Hsapiens.dbSNP132.db* and *PolyPhen.Hapiens.dbSNP131.db* and are designed to be searched by `rsid`. Variants that are in dbSNP can be searched with these database packages. When working with novel variants, SIFT and PolyPhen must be called directly. See references for home pages.

The pre-calculated predictions from SIFT and PolyPhen are based on particular gene models. SIFT uses Ensembl and PolyPhen uses the UCSC Known Genes track. It is important that the annotation file used to identify coding / non-coding variants is based on the same gene model as these predictions. We will be using SIFT and the TranscriptDb we used had Ensembl gene ids.

Identify the non-synonymous variants and obtain the rsids.

```
> nms <- names(coding)
> idx <- values(coding)$CONSEQUENCE == "nonsynonymous"
> nonsyn <- coding[idx]
> names(nonsyn) <- nms[idx]
> rsids <- unique(names(nonsyn)[grep("rs", names(nonsyn), fixed=TRUE)])
```

Detailed descriptions of the database columns can be found with `?SIFTdbColumns` and `?PolyPhenDbColumns`. Variants in these databases often contain more than one row per variant. The variant may have been reported by multiple sources and therefore the source will differ as well as some of the other variables.

```
> library(SIFT.Hsapiens.dbSNP132)
> ## rsids in the package
> head(keys(SIFT.Hsapiens.dbSNP132), 3)

[1] "rs10000692" "rs10001580" "rs10002700"

> ## list available columns
> cols(SIFT.Hsapiens.dbSNP132)

[1] "RSID"          "PROTEINID"      "AACHANGE"       "METHOD"         "AA"
[6] "PREDICTION"    "SCORE"          "MEDIAN"         "POSTIONSEQS"    "TOTALSEQS"

> ## select a subset of columns
> ## a warning is thrown when a key is not found in the database
> subst <- c("RSID", "PREDICTION", "SCORE", "AACHANGE", "PROTEINID")
> sift <- select(SIFT.Hsapiens.dbSNP132, keys=rsids, cols=subst)
> head(sift, 3)
```

	RSID	PROTEINID	AACHANGE	PREDICTION	SCORE
1	rs114335781	<NA>	<NA>	<NA>	<NA>
2	rs8135963	<NA>	<NA>	<NA>	<NA>
3	rs114264124	NP_077010	R233Q	TOLERATED	0.59

PolyPhen provides predictions using two different training datasets and has considerable information about 3D protein structure. See `?PolyPhenDbColumns` or the PolyPhen web site listed in the references for more details.